

Creation Date: 24/01/2020 18:12

Last Modified Date: 27/01/2020 17:19

Lec 13 : Generative Models

Lecture 13: Generative Models

Fei-Fei Li & Justin Johnson & Serena Yeung

Lecture 13 - 1

University
May 18, 2017

Overview

- Unsupervised Learning
- Generative Models
 - Restricted Boltzmann Machine (RBM) and Deep Boltzmann Machine (DBM)

- Generative Models: It is a type of Unsupervised Learning.

Supervised vs Unsupervised Learning

Supervised Learning



Supervised vs Unsupervised Learning

Supervised Learning



Supervised vs Unsupervised Learning

Supervised Learning



Supervised vs Unsupervised Learning

Supervised Learning



- Goal of Unsupervised Learning: To learn the underlying structure of the given data.

Supervised vs Unsupervised Learning

Supervised vs Unsupervised Learning

Supervised vs Unsupervised Learning

Supervised vs Unsupervised Learning

- Density Estimation: Modeling the data such that the probability of the data generation is high where more points are concentrated.

- Differences between Supervised and Unsupervised Learning:

Supervised vs Unsupervised Learning

- Solving Unsupervised Learning problems helps in understanding the structure of the visual world.

Generative Models

- **Generative models address density estimation.**
- **Explicit Density Estimation:** Explicitly define and solve for $p_{model}(x)$
- **Implicit Density Estimation:** We learn a model that can sample from $p_{model}(x)$ **without** explicitly defining it.

Why Generative Models?

Taxonomy of Generative Models

Direct

PixelRNN and PixelCNN:

PixelRNN and PixelCNN

Fully visible belief network

PixelRNN [van der Oord et al. 2016]

PixelRNN is a recurrent neural network that processes images pixel by pixel.

PixelRNN is a recurrent neural network that processes images pixel by pixel.

PixelCNN [van der Oord et al. 2016]

PixelCNN is a convolutional neural network that processes images pixel by pixel.

PixelCNN is a convolutional neural network that processes images pixel by pixel.

PixelCNN *[van der Oord et al. 2016]*

Generating images from a single pixel

- The generation depends on the value of the starting pixel, because everything else is conditioned on that initial pixel.

Generation Samples

PixelRNN and PixelCNN

Variational Autoencoders

Variational

Autoencoders Variational Autoencoders Generative Adversarial Networks

So far...

DiverGNNs define tractable density function, optimize likelihood of training data.

Some background first: Autoencoders

- Ans: We want dimensionality reduction since we want a compact representation of x but also which represent all the important features of x .

Some background first: Autoencoders

Some background first: Autoencoders Reconstructed data

Some background first: Autoencoders Reconstructed data

Some background first: Autoencoders

Some background first: Autoencoders

- ***So using large amount of unlabeled data we can learn a better feature representation and then use it to initialize a supervised learning problem where we only have small data.***
- This approach helps because with small amount of data supervised approaches will cause overfitting.
- So using this approach we get a better initialization of the supervised model with better features.
- ***Features capture factors of variation in training data.***

Some background first: Autoencoders

Variational Autoencoders

- Variational Autoencoders help us in sampling
- The intuition is that each element $z^{(i)}$ captures certain feature of x (like smile, eye color of faces).
- z is modeling using some priors, more commonly we use Gaussian priors (μ, σ) .

Variational Autoencoders

We want to estimate the true parameters θ^*
of this generative model

- We choose the prior $p(z)$ to simple, like Gaussian. Because it is reasonable for latent attributes.
- The conditional $p(x|z)$ is complex in nature since it is generating images where $p(z)$ is just a vector so we try to model z using Gaussian.

Variational Autoencoders

We want to estimate the true parameters θ^*
of this generative model

Variational Autoencoders

We want to estimate the true parameters θ^*
of this generative model

Variational Autoencoders

We want to estimate the true parameters θ^* of this generative model

Variational Autoencoders: Intractability

Data likelihood: $p_\theta(x) = \int p_\theta(z) p_\theta(x|z) dz$ ✓

Variational Autoencoders: Intractability

Data likelihood: $p_z(x) = \int p_z(z) p_x(x|z) dz$

Variational Autoencoders: Intractability

Data likelihood: $p_z(x) = \int p_z(z) p_x(x|z) dz$

- The integral is intractable.

Variational Autoencoders: Intractability

Data Likelihood: $p_+(x) = \int p_+(x) p_-(x|z) dz$

Variational Autoencoders: Intractability

Data Likelihood: $p_+(x) = \int p_+(x) p_-(x|z) dz$

Variational Autoencoders

Since we're modeling probabilistic generation of data, encoder and decoder networks are probabilistic.

Variational Autoencoders

Since we're modeling probabilistic generation of data, encoder and decoder networks are probabilistic.

Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood.

Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood.

Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood.

Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood.

Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood.

Variational Autoencoders

Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood.

Variational Autoencoders

Putting it all together: maximizing the

Variational Autoencoders

Putting it all together: maximizing the

Variational Autoencoders

Putting it all together: maximizing the

Variational Autoencoders

Putting it all together: maximizing the

Variational Autoencoders

Put this all together to maximize the

Maximize

\hat{x}

Variational Autoencoders

Put this all together to maximize the

Maximize

\hat{x}

Variational Autoencoders: Generating Data!

Use decoder network. Draw sample z from prior.

Variational Autoencoders: Generating Data!

Use decoder network. Draw sample z from prior. Determine if z is good.

Variational Autoencoders: Generating Data!



Variational Autoencoders: Generating Data!



Variational Autoencoders

Probabilistic spin to traditional autoencoders => allows generating data

GANs:

So far...

PixelCNNs define tractable density function, optimize likelihood of training data:

$$p(x) = \prod_{i=1}^n p(x_i | x_{1:i-1})$$

Generative Adversarial Networks

Ian Goodfellow et al., "Generative Adversarial Nets", NIPS 2014

Training GANs: Two-player game

Ian Goodfellow et al., "Generative
Adversarial Nets", NIPS 2014

Training GANs: Two-player game

Ian Goodfellow et al., "Generative
Adversarial Nets", NIPS 2014

Training GANs: Two-player game

Ian Goodfellow et al., "Generative
Adversarial Nets", NIPS 2014

Training GANs: Two-player game

Ian Goodfellow et al., "Generative
Adversarial Nets", NIPS 2014

Training GANs: Two-player game

Ian Goodfellow et al., "Generative
Adversarial Nets", NIPS 2014

Generative Adversarial Nets

Ian Goodfellow et al., "Generative
Adversarial Nets", NIPS 2014

Generative Adversarial Nets

Ian Goodfellow et al., "Generative Adversarial Nets", NIPS 2014


Generative Adversarial Nets (GANs)

Generative Adversarial Nets: Convolutional Architectures

Generative Adversarial Nets (GANs) with Convolutional Architectures

Generative Adversarial Nets: Convolutional Architectures

128



Generative Adversarial Nets: Convolutional Architectures

Generative Adversarial Nets: Convolutional Architectures

Generative Adversarial Nets: Interpretable Vector Math

Radford et al. ICLR 2016

Generative Adversarial Nets: Interpretable Vector Math

Christopher Manning, Andrew Senior, Alex Senior, David Rosenberg, and Ilya Sutskever

Radford et al

2017: Year of the GAN

Text -> Image Synthesis

this small bird has a pink this magnificent fellow is
bracer and crown and black almost all black with a red

“The GAN Zoo”

See also: <https://github.com/soumith/ganhacks> for tips and tricks for trainings GANs

GANs

Don't work with an explicit density function

Take game-theoretic approach: learn to generate from training distribution through 2-player

Recap

Generative Models

- **DiagonalPNN and DiagonalCNN** Explicit density model, optimizes exact likelihood, good